

# 1. Understand the networking configuration on the cluster nodes

Reference is [here](#).

There are 4 distinct networking problems to solve:

- Highly-coupled container-to-container communications
- Pod-to-Pod communications
- Pod-to-Service communications
- External-to-internal communications

## Explore Kubernetes Environment

- What is the port the kube-scheduler is listening on in the master node?
- Notice that ETCD is listening on two ports. Which of these have more client connections established?
- Inspect the kubelet service and identify the network plugin configured for Kubernetes.
- What is the path configured with all binaries of CNI supported plugins?
- What is the CNI plugin configured to be used on this kubernetes cluster?

show

```
netstat -nplt
netstat -anp | grep etcd
netstat -anp | grep etcd | grep 2379 | wc -l
netstat -anp | grep etcd | grep 2380 | wc -l
ps -aux | grep kubelet
systemctl status kubelet.service
ls /etc/cni/net.d/
```

## Choosing different CNI networking solutions

- Deploy weave-net networking solution to the cluster
- Deploy calico networking solution to the cluster
- Deploy flannel networking solution to the cluster

show

- Deploy weave-net. Documentation is [here](#).

```
kubect1 apply -f "https://cloud.weave.works/k8s/net?k8s-version=$(kubect1 version | base64 | tr -d '\n')"
```

- Deploy calico. Documentation is [here](#).

```
kubect1 apply -f https://docs.projectcalico.org/v3.8/manifests/calico.yamlhttps://docs.projectcalico.org/v3.8/manifests/calico.yaml
```

- Deploy flannel. Documentation is [here](#).

```
kubect1 apply -f https://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.ymlhttps://raw.githubusercontent.com/coreos/flannel/master/Documentation/kube-flannel.yml
```

# 2. Understand Pod networking concepts

show

# 3. Understand Service Networking

## Service networking

- What network range are the nodes in the cluster part of? Run the command `ip addr` and look at the IP address assigned to the `ens3` interfaces. Derive network range from that.
- What is the range of IP addresses configured for PODs on this cluster?
  - For weave-net, Check the weave pods logs using command `kubectl logs weave -n kube-system` and look for `ipalloc-range`
  - For calico, check the `calico.yaml` file and looking for `CALICO_IPV4POOL_CIDR`
- What is the IP Range configured for the services within the cluster?
  - Inspect the setting on kube-api server by running on command `ps -aux | grep kube-api`
- What type of proxy is the kube-proxy configured to use?
  - Check the logs of the kube-proxy pods. Command: `kubectl logs kube-proxy-ft6n7 -n kube-system`

show

## 4. Deploy and configure network load balancer

show

## 5. Know how to use Ingress rules

Note that the innotation in setting the Ingress Controller in very **IMPORTANT**: Reference is [here](#). We can select between `gce/nginx/traefik`.

For instance,

```
metadata:
  name: foo
  annotations:
    kubernetes.io/ingress.class: "gce"
```

will target the GCE controller, forcing the nginx controller to ignore it, while an annotation like

```
metadata:
  name: foo
  annotations:
    kubernetes.io/ingress.class: "nginx"
```

will target the nginx controller, forcing the GCE controller to ignore it.

### Traefik ingress controller

Installation guide is [here](#).

- Create a ClusterRole 'traefik-ingress-controller' with two rules:
  - `apiGroup: ""`, resources: services, endpoints, secrets, verbs: get, list, watch
  - `apiGroup: extensions`, resources: ingresses, verbs: get, list, watch
- Create a ClusterRoleBiding 'traefik-ingress-controller' that binds the ClusterRole 'traefik-ingress-controller' and the service account 'traefik-ingress-controller'.
- Deploy Traefik using the DaemonSet:
  - Download and save the ds file using [this link](#).
  - Edit the file, delete securityContext lines and add 'hostNetwork: true' line that lines up with containers: line.
  - Create the ingress controller with `kubectl create -f`.
- Create a deployment name 'secondapp', with image 'nginx', then expose the new server as NodePort, port 80.
- Create an Ingress 'ingress-test', with the rules:
  - Host: [www.secondapp.com](http://www.secondapp.com)
  - Backend: secondapp service, port 80
- Check with `curl -H "Host: www.secondapp.com" http://20.0.0.11/`
- Create a deployment name 'thirdpage', with image 'nginx', then expose the new server as NodePort, port 80.
- Execute the thirdpage pod to modify the title of the webpage to 'Third Page' locate at `/usr/share/nginx/html/index.html`
- Modify the ingress-test, adding second rules:
  - Host: [www.thirdpage.com](http://www.thirdpage.com)

- Check Backend thirdpage: <http://20.0.0.11/>

show

- Create a ClusterRole 'traefik-ingress-controller'

```
kubectl create clusterrole traefik-ingress-controller --resource=ingresses --verb=get,list,watch \
--dry-run -o yaml > traefik-cluster-role.yaml
vim traefik-cluster-role.yaml
kubectl create -f traefik-cluster-role.yaml
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: traefik-ingress-controller
rules:
- apiGroups:
  - extensions
  resources:
  - ingresses
  verbs:
  - get
  - list
  - watch
- apiGroups:
  - ""
  resources:
  - services
  - endpoints
  - secrets
  verbs:
  - get
  - list
  - watch
```

- Create a ClusterRoleBinding 'traefik-ingress-controller'

```
kubectl create clusterrolebinding traefik-ingress-controller --clusterrole=traefik-ingress-controller \
--serviceaccount=kube-system:traefik-ingress-controller
```

- Deploy Traefik using the DaemonSet:

```
wget https://raw.githubusercontent.com/containous/traefik/v1.7/examples/k8s/traefik-ds.yamlhttps://raw.githubusercontent.com/containous/traefik/v1.7/examples/k8s/traefik-ds.yaml -O traefik-ds.yaml
vim traefik-ds.yaml
kubectl create -f traefik-ds.yaml
```

- Create deployment secondapp

```
kubectl create deployment secondapp --image=nginx
kubectl expose deployment secondapp --type=NodePort --port=80
```

- Create an Ingress 'ingress-test': reference is [here](#).

```
vim ingress-test.yaml
```

```
apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: ingress-test
spec:
  rules:
  - host: www.secondapp.comwww.secondapp.com
    http:
      paths:
      - backend:
          serviceName: secondapp
          servicePort: 80
```

```
kubectl create -f ingress-test.yaml
curl -H "Host: www.secondapp.com"www.secondapp.com http://20.0.0.11http://20.0.0.11
```

- Create deployment thirdpage, edit the deployment to add the [custom dnspolicy](#).

```
kubectl create deployment thirdpage --image=nginx
```

```
kubectrl expose deployment thirdpage --type=NodePort --port=80
kubectrl edit deployment thirdpage
```

```
spec:
  containers:
  - image: nginx
    imagePullPolicy: Always
    name: nginx
  dnsConfig:
    nameservers:
    - 8.8.8.8
  dnsPolicy: None
```

- Execute the thirdpage pod to modify nginx webpage title.

```
kubectrl exec thirdpage-658458994f-77mhs -it -- /bin/bash
apt-get update
apt-get install vim -y
vim /usr/share/nginx/html/index.html
```

- Modify the ingress-test, adding second rules:

```
kubectrl edit ingress ingress-test

spec:
  rules:
  - host: www.secondapp.comwww.secondapp.com
    http:
      paths:
      - backend:
          serviceName: secondapp
          servicePort: 80
  - host: www.thirdpage.comwww.thirdpage.com
    http:
      paths:
      - backend:
          serviceName: thirdpage
          servicePort: 80
```

- Check with curl

```
curl -H "Host: www.thirdpage.com"www.thirdpage.com" http://20.0.0.11http://20.0.0.11
curl -H "Host: www.secondapp.com"www.secondapp.com" http://20.0.0.11http://20.0.0.11
```

## Nginx ingress controller

- Create 2 new namespaces 'app-space', 'critical-space'
- Create 4 deployments:
  - default-backend (image: kodekloud/ecommerce:404),
  - webapp-food (image: kodekloud/ecommerce:food),
  - webapp-video (image: kodekloud/ecommerce:video),
  - webapp-wear (image: kodekloud/ecommerce:apparels).
- Expose 4 deployments of type ClusterIP:
  - default-http-backend: port 80
  - food-service: port 8080
  - video-service: port 8080
  - wear-service: port 8080
- Create ingress name 'ingress-wear-watch' in 'app-space', to have the service at different paths: /wear, /stream, /eat.
- Create deployment: webapp-pay (image: kodekloud/ecommerce:pay) in 'critical-space', expose as 'pay-service' at port 8282.
- Create ingress name 'ingress-pay' in 'critical-space', to have the service at path: /pay.

show

- Create two namespaces:

```
kubectrl create ns app-space
kubectrl create ns critical-space
```

- Create 4 deployments

```
kubectl create deployment default-backend --image=kodekloud/ecommerce:404 --namespace=app-space
kubectl create deployment webapp-food --image=kodekloud/ecommerce:food --namespace=app-space
kubectl create deployment webapp-video --image=kodekloud/ecommerce:video --namespace=app-space
kubectl create deployment webapp-wear --image=kodekloud/ecommerce:apparels --namespace=app-space
```

- Expose 4 deployments:

```
kubectl -n app-space expose deployment default-backend --type=ClusterIP --port=80 --name=default-http-backend
kubectl -n app-space expose deployment webapp-food --type=ClusterIP --port=8080 --name=food-service
kubectl -n app-space expose deployment webapp-video --type=ClusterIP --port=8080 --name=video-service
kubectl -n app-space expose deployment webapp-wear --type=ClusterIP --port=8080 --name=wear-service
```

- Create ingress name 'ingress-wear-watch': **ANNOTATIONS** is very **IMPORTANT**. Reference is [here](#).

```
vim ingress-wear-watch.yaml

apiVersion: networking.k8s.io/v1beta1
kind: Ingress
metadata:
  name: ingress-wear-watch
  namespace: app-space
  annotations:
    kubernetes.io/ingress.class: traefik
    traefik.frontend.rule.type: PathPrefixStrip
spec:
  rules:
  - http:
      paths:
      - path: /wear
        backend:
          serviceName: wear-service
          servicePort: 8080
      - path: /stream
        backend:
          serviceName: video-service
          servicePort: 8080
      - path: /eat
        backend:
          serviceName: food-service
          servicePort: 8080

kubectl create -f ingress-wear-watch.yaml
```

## 6. Know how to configure and use the cluster DNS

### Explore CoreDns

show

## 7. Understand CNI

show