

****Lab Activity 1 - Label Selectors****

Deploy two pods:

One pod with a label of “Tier = Web”

```
kubect1 run web --image=nginx --labels="Tier=Web"
```

Or

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    Tier: App
  name: app
spec:
  containers:
    - image: nginx
      name: app
```

One pod with a label of “Tier = App”

```
kubect1 run app --image=nginx --labels="Tier=App"
```

Or

```
apiVersion: v1
kind: Pod
metadata:
  labels:
    Tier: Web
  name: web
spec:
  containers:
    - image: nginx
      name: web
```

Verify the labels are applied.

```
kubect1 get pods --show-labels
```

| NAME | | READY | STATUS | RESTARTS | AGE | LABELS |
|------|-----|---------|--------|----------|-----|----------|
| app | 1/1 | Running | 0 | 24s | | Tier=App |
| Web | 1/1 | Running | 0 | 71s | | Tier=Web |

****Lab Activity 2 - Daemonsets****

Deploy a Daemonset that leverages the nginx image

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: nginx-ds
spec:
  selector:
    matchLabels:
      name: nginx-ds
  template:
    metadata:
      labels:
        name: nginx-ds
    spec:
      containers:
        - name: nginx-ds
          image: nginx
```

Verify the daemonset has been created successfully

```
kubect1 get daemonset
NAME          DESIRED  CURRENT  READY  UP-TO-DATE  AVAILABLE  NODE SELECTOR  AGE
```

****Lab Activity 3 - Resource Limits****

Create a new namespace called “tenant-b-100mi”

```
kubectl create ns tenant-b-100mi
namespace/tenant-b-100mi created
```

Create a memory limit of 100Mi for this namespace

```
apiVersion: v1
kind: LimitRange
metadata:
  name: mem-limit-range
spec:
  limits:
  - max:
      memory: 100Mi
    type: Container
```

Create a pod with a memory request of 150Mi, ensure the limit has been set by verifying you

get a error message.

```
apiVersion: v1
kind: Pod
metadata:
  name: default-mem-demo
spec:
  containers:
  - name: default-mem-demo-ctr
    image: nginx
    resources:
      requests:
        memory: 150Mi
```

The Pod "default-mem-demo" is invalid: spec.containers[0].resources.requests: Invalid value: "150Mi": must be less than or equal to memory limit

****Lab Activity 4 - Multiple Schedulers****

Assume another scheduler “custom-scheduler” has been created in your environment. Configure a pod to use this scheduler.

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-web
  labels:
    role: web
spec:
  containers:
  - name: nginx
  image: nginx
  schedulerName: custom-scheduler
```

Validate the pod is using this scheduler.

```
kubectl get pod nginx-web -o yaml | grep schedulerName
```

```
{"apiVersion":"v1","kind":"Pod","metadata":{"annotations":{},"labels":{"role":"web"},"name":"nginx-web","namespace":"default"},"spec":{"containers":[{"image":"nginx","name":"nginx"}],"schedulerName":"custom-scheduler"}}
```

```
schedulerName: custom-scheduler
```

****Lab Activity 5 - Schedule Pod without a scheduler****

One one of the worker nodes:

Create a the directory /etc/staticpods

```
mkdir /etc/staticpods
```

Create a pod manifest file in this directory

```
apiVersion: v1
kind: Pod
metadata:
  name: staticpod
spec:
  containers:
  - name: staticpod
    image: nginx
```

Configure the kubelet service on this worker node to create pods from /etc/staticpods

```
● kubelet.service - kubelet: The Kubernetes Node Agent
   Loaded: loaded (/lib/systemd/system/kubelet.service; enabled; vendor preset: enabled)
   Drop-In: /etc/systemd/system/kubelet.service.d
            └─10-kubeadm.conf
   Active: active (running) since Mon 2019-04-29 12:07:52 UTC; 2min 50s ago
     Docs: https://kubernetes.io/docs/home/https://kubernetes.io/docs/home/
  Main PID: 8099 (kubelet)
    Tasks: 16 (limit: 2320)
   CGroup: /system.slice/kubelet.service
            └─8099 /usr/bin/kubelet --bootstrap-kubeconfig=/etc/kubernetes/bootstrap-kubelet.conf --kubeconfig=/etc/kubernetes/kubelet.conf --config=/var/lib/kubelet/config.yaml
```

```
Nano /var/lib/kubelet/config.yaml
```

Add the following line

```
staticPodPath: /etc/staticpods
```

```
sudo systemctl daemon-reload
sudo systemctl restart kubelet
```

****Lab Activity 6 - Display Scheduler Events****

Create a pod manifest file using the nginx image which will create a pod called “nginx-web” (Alternatively do this via kubectl run)

```
kubectl run nginx --image=nginx
```

Extract the events from the cluster, particularly those pertaining to scheduling to find where this pod was scheduled to.

```
kubectl describe pod nginx-7db9fccd9b-xw6qd
```

Extract the logs from the pod running the default scheduler, or from the respective file if running as a deamon service on your master node.

```
kubectl get pods -n kube-system | grep sch
kube-scheduler-k8s-master-03      1/1   Running    0      3h1m
```

```
kubectl logs kube-scheduler-k8s-master-03 -n kube-system
```

****Lab Activity 7 - Know how to configure the Kubernetes Scheduler****

Configure the Kube-Scheduler by adding --logtostderr=true to the existing configuration.

```
sudo cat /etc/kubernetes/manifests/kube-scheduler.yaml
```

```
spec:
  containers:
  - command:
  - kube-scheduler
  - --bind-address=127.0.0.1
  - --kubeconfig=/etc/kubernetes/scheduler.conf
  - --leader-elect=true
  - --logtostderr=true
```

****Lab Activity 8 - Taints****

Add taint a node

```
kubect1 taint nodes node-01 available=no:NoExecute
```

Pod manifest with tolerations

```
apiVersion: v1
kind: Pod
metadata:
  name: activity-8
spec:
  containers:
  - name: activity-8-ctr
    image: nginx
  tolerations:
  - key: "available"
    operator: "Equal"
    value: "no"
    effect: "NoExecute"
```

Observe behaviour by applying manifest and using `kubect1 get pods -o wide` to view placement of new pod or existance of existing pod.

Pod manifest with tolerationSeconds

```
apiVersion: v1
kind: Pod
metadata:
  name: activity-8
spec:
  containers:
  - name: activity-8-ctr
    image: nginx
  tolerations:
  - key: "key"
    operator: "Equal"
    value: "value"
    effect: "NoExecute"
    tolerationSeconds: 15
```